

Predict whether the client has subscribed a term deposit or not.

Import necessary libraries

```
In [1]: import warnings
```

```
In [2]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
from matplotlib import style

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import LabelEncoder, StandardScaler, OrdinalEncoder

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [3]: # Matplotlib configurations

# Display interactive plots. Used this since convenient for displaying
# %matplotlib notebook
# Font and figure size:
# Ref: https://stackoverflow.com/questions/3899980/how-to-change-the-font-size
SMALL_SIZE = 8
MEDIUM_SIZE = 9
BIGGER_SIZE = 12

plt.rc('font', size=SMALL_SIZE)          # controls default text sizes
plt.rc('axes', titlesize=SMALL_SIZE)     # fontsize of the axes title
plt.rc('axes', labelsize=MEDIUM_SIZE)    # fontsize of the x and y labels
plt.rc('xtick', labelsize=SMALL_SIZE)     # fontsize of the tick labels
plt.rc('ytick', labelsize=SMALL_SIZE)     # fontsize of the tick labels
plt.rc('legend', fontsize=SMALL_SIZE)     # legend fontsize
```

Import dataset

```
In [4]:
```

```
In [5]:
```

```
Out[5]:
```

| | age | job | marital | education | default | balance | housing | loan | contact | day | mon |
|---|-----|--------------|---------|-----------|---------|---------|---------|------|---------|-----|-----|
| 0 | 58 | management | married | tertiary | no | 2143 | yes | no | unknown | 5 | m |
| 1 | 44 | technician | single | secondary | no | 29 | yes | no | unknown | 5 | m |
| 2 | 33 | entrepreneur | married | secondary | no | 2 | yes | yes | unknown | 5 | m |

| | age | job | marital | education | default | balance | housing | loan | contact | day | mon |
|---|-----|-------------|---------|-----------|---------|---------|---------|------|---------|-----|-----|
| 3 | 47 | blue-collar | married | unknown | no | 1506 | yes | no | unknown | 5 | m |

EDA

In [6]:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
#   Column          Non-Null Count  Dtype
---  -
0   age              45211 non-null  int64
1   job              45211 non-null  object
2   marital          45211 non-null  object
3   education        45211 non-null  object
4   default          45211 non-null  object
5   balance          45211 non-null  int64
6   housing          45211 non-null  object
7   loan             45211 non-null  object
8   contact          45211 non-null  object
9   day              45211 non-null  int64
10  month            45211 non-null  object
11  duration         45211 non-null  int64
12  campaign         45211 non-null  int64
13  pdays            45211 non-null  int64
14  previous         45211 non-null  int64
15  poutcome        45211 non-null  object
16  y                45211 non-null  object
dtypes: int64(7), object(10)
memory usage: 5.9+ MB
```

Observations:

- Data is from a campaign conducted by the bank, to get their customers to make term deposits.
- Data contains information regarding the banks's customers i.e their personal details, bank related details, campaign details and corresponding outcome; whether the customer made a term deposit. It is a mix of numeric and categorical variables.
- There are six different kinds of variables in this dataset
 - Cyclic (numeric and cat) - "day" and "month"
 - Interval (numeric) - "age", "duration", "campaign", "pdays"
 - Categorical; ordinal - "education"
 - Categorical; Label; binary - "default", "housing", "loan", "y"
 - Categorical; nominal - "job", "marital", "contact", "poutcome"
 - Numeric - "previous", "balance"
- We must use different encoding techniques in order to properly process the dataset.
- There are no null values.

```
In [7]: # Extracting column names and sorting them to appropriate categories.
n_cols = len(bank_df.columns)
numeric_cols = []
cat_cols = []
```

```

for i in range(n_cols):
    if bank_df.dtypes[i] == 'int64':
        numeric_cols.append(bank_df.columns[i])
    else:
        cat_cols.append(bank_df.columns[i])

X_cols = [col for col in bank_df.columns if col not in ("y")]
y_col = "y"

print("Numeric columns      :", numeric_cols)
print("Categorical columns  :", cat_cols)

cat_cols.remove("y")

print("Categorical features :", cat_cols)
print("Features             :", X_cols)

Numeric columns      : ['age', 'balance', 'day', 'duration', 'campaign', 'pdays', 'previous']
Categorical columns  : ['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'outcome', 'y']
Categorical features : ['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'outcome']
Features             : ['age', 'job', 'marital', 'education', 'default', 'balance', 'housing', 'loan', 'contact', 'day', 'month', 'duration', 'campaign', 'pdays', 'previous', 'outcome']
Target               : y

```

Model building

```

In [8]: # Separating features and target.
X = bank_df.drop(['y'], axis=1)

```

```

In [9]: # Splitting the data into training and testing sets.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

```

```

Out[9]: ((30291, 16), (14920, 16), (30291,), (14920,))

```

```

In [10]: # Encode categorical features.
ordinal_enc = OrdinalEncoder()
X_train_enc = pd.DataFrame(ordinal_enc.fit_transform(X_train[cat_cols]))
X_test_enc = pd.DataFrame(ordinal_enc.transform(X_test[cat_cols]))

# Encoding removes column names, thus we need to insert them back
X_train_enc.columns = X_train[cat_cols].columns

```

```

In [11]: # Dropping corresponding categorical columns.
X_train.drop(cat_cols, axis=1, inplace=True)
X_test.drop(cat_cols, axis=1, inplace=True)

# Concatenating the encoded columns with the original data.
X_train = pd.concat([X_train, X_train_enc.set_index(X_train.index)], axis=1)

```

```

In [12]: # Encoding the target variable.
label_enc = LabelEncoder()
y_train_enc = label_enc.fit_transform(y_train)
y_test_enc = label_enc.transform(y_test)

```

```
In [13]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
```

Model training

```
In [14]: log_model = LogisticRegression()
```

```
Out[14]: LogisticRegression()
```

```
In [15]:
```

```
Out[15]: array([[ 0.07149251,  0.06902964, -0.03881063,  1.02909677, -0.39319
337,
               0.36962832,  0.18706114,  0.02638061,  0.13285927,  0.14560
422,
              -0.05506589, -0.50317089, -0.24899631, -0.56736305,  0.11462
209,
               0.22960332]])
```

```
In [16]:
```

```
Out[16]: array([-2.68050071])
```

Model testing

```
In [17]: y_pred_train = log_model.predict(X_train)
```

Predictions and evaluation for train data

```
In [18]:
```

```
Accuracy score: 0.8905945660427189
```

```
In [19]:
```

```
Confusion matrix:
[[26218  529]
 [ 2785  759]]
```

```
In [20]:
```

```
Precision score: 0.5892857142857143
```

```
In [21]:
```

```
Recall score: 0.2141647855530474
```

```
In [22]:
```

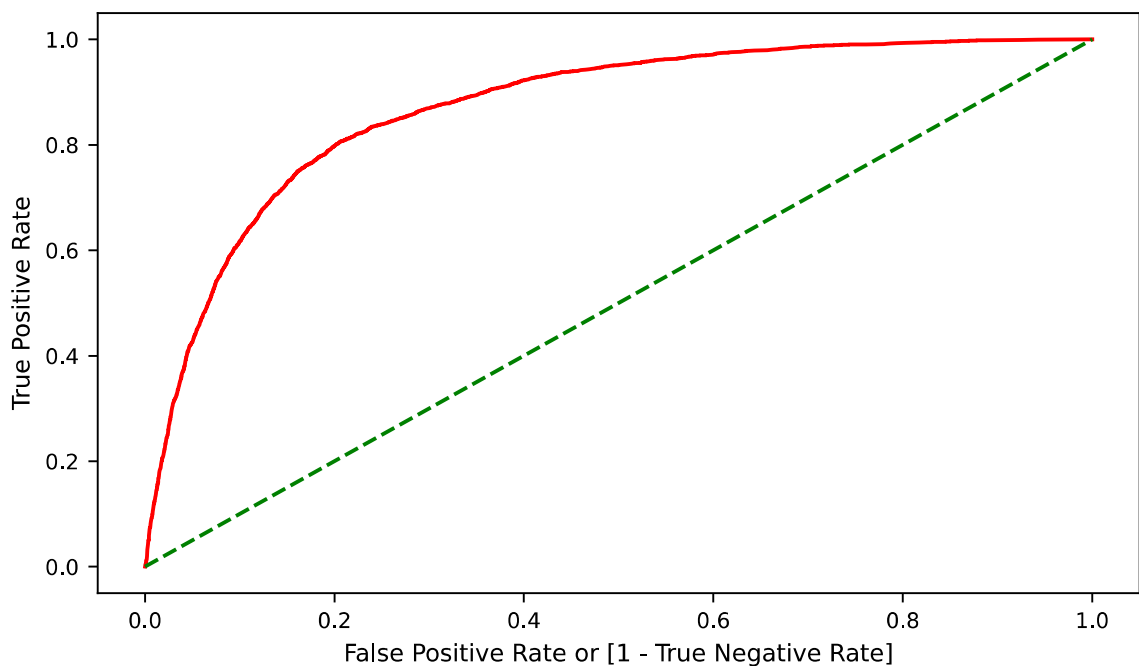
| | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 0.90 | 0.98 | 0.94 | 26747 |
| 1 | 0.59 | 0.21 | 0.31 | 3544 |

ROC curve - train data

```
In [23]: fpr, tpr, thresholds = roc_curve(y_train_enc, log_model.predict_proba (
auc = roc_auc_score(y_train_enc,y_pred_train)
print(auc)

fig,ax = plt.subplots()
ax.plot(fpr, tpr, color='red', label='logit model ( area = %0.2f)'%auc
ax.plot([0, 1], [0, 1], 'g--')
ax.set_xlabel('False Positive Rate or [1 - True Negative Rate]')
ax.set_ylabel('True Positive Rate')
print(auc)
```

0.597193433267046



Predictions and evaluation for test data

```
In [24]:
```

Accuracy score: 0.8916890080428954

```
In [25]:
```

Confusion matrix:

```
[[12931  244]
 [ 1372  373]]
```

For this problem, although the number of false negatives are more(as compared to false positives), it is good in a way since, bank could launch a much more vigorous campaign next

year so that more customers would make term deposit. Also, the number of false positives is very less compared to true positives. Although this is the case there is a huge difference between precision and recall scores for 1 class. For a good model, this difference should be within 10%. Thus improvement has to be made with respect to getting the scores as close as possible. Since it is doing a good job predicting the 0 class, we can still use the model but it

In [26]:

```
Precision score: 0.6045380875202593
```

In [27]:

```
Recall score: 0.21375358166189112
```

In [28]:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.90 | 0.98 | 0.94 | 13175 |
| 1 | 0.60 | 0.21 | 0.32 | 1745 |
| accuracy | | | 0.89 | 14920 |
| macro avg | 0.75 | 0.60 | 0.63 | 14920 |
| weighted avg | 0.87 | 0.89 | 0.87 | 14920 |


ROC curve -test data:

```
In [29]: fpr, tpr, thresholds = roc_curve(y_test_enc, log_model.predict_proba (X

auc = roc_auc_score(y_test_enc,y_pred_test)
print(auc)

fig,ax = plt.subplots()
ax.plot(fpr, tpr, color='red', label='logit model ( area = %0.2f)'%auc
ax.plot([0, 1], [0, 1], 'g--')
ax.set_xlabel('False Positive Rate or [1 - True Negative Rate]')
ax.set_ylabel('True Positive Rate')

0.5976168287816097
```



Conclusion:

A logistic regression model was constructed to predict whether the client has subscribed to a term deposit or not.

Type *Markdown* and LaTeX: α^2